

# AI Agents and Search

---

Lara J. Martin (she/they)

<https://laramartin.net/neurosymbolic-text-gen>

*Slides adapted from Chris Callison-Burch, Cynthia Matuszek, & Cassandra Kent*

# Learning Objectives

---

Understand the difference between traditional AI agents and agentic AI

Categorize agents based on their capabilities and behavior

Formulate a search problem

Review basic types of tree search algorithms

# AI Agents

---

# AI Agent Definition


**Agent:** anything that **perceives** its environment through **sensors**, and **acts** on its environment through **actuators**

**Percept:** input at an instant

**Percept sequence:** history of inputs

**Agent function:** mapping of **percept sequence** to **action**

**Agent program:** (concise) implementation of an agent function



*Picture from Dr. Cynthia Matuszek, source unknown*

# “Agentic AI”

---

Term coined around 2023 but has become popular as of early 2025

Using a *language model* to behave like an “agent” – the intersection of AI & NLP beyond chatbots

“autonomous systems designed to pursue complex goals with *minimal human intervention*” [1]

But technically all AI is agentic...

# Types of Agents


---

Similar to environments, we can come up with a taxonomy of agents based on how they behave:

1. Simple Reflex Agent
2. Model-based Reflex Agent
3. Model-based Goal-based Agent
4. Model-based Utility-based Agent

# Example: Vampire World

---




# Types of Agents:

## Simple Reflex Agent

Decide actions directly from the current percept sequence

Example:

```
if [A, Empty]: return Right
if [B, Empty]: return Left
if [A or B, Human]: return Suck
```



*Artificial Intelligence: A Modern Approach, Pg. 50*

# Types of Agents:


## Model-based Reflex Agent

Attempt to model the state of the world, decide action based on modeled (“remembered”) state

Example:

Observe that at a certain time, people are more likely to show up in A

Move to location based on time of day



*Artificial Intelligence: A Modern Approach, Pg. 52*

# Types of Agents:


## Model-based Goal-based Agent

Action choices change as the goal changes

Example:

For the same performance criteria (suck as much blood as possible), we can add a goal of keeping one room clear to act as a home

As the goal changes, vampire changes how often it moves



*Artificial Intelligence: A Modern Approach, Pg. 53*

# Types of Agents:

## Model-based Utility-based Agent


Agent cares not only about reaching a goal, but also *how* it reaches its goal

Example:

Energy-efficient vampire agent

Still wants to suck as much blood as possible

Limits how often it moves



*Artificial Intelligence: A Modern Approach, Pg. 55*

# What type of agent is a Roomba?

---



1. Simple Reflex Agent
- ✓ 2. Model-based Reflex Agent ← Keeps a memory of the world but only reacts to its surroundings
3. Model-based Goal-based Agent
4. Model-based Utility-based Agent

# What type of agent is Watson?



<https://cdn.geekwire.com/wp-content/uploads/2017/04/IBM-Watson.jpg>

Watson was an agent made by IBM that played *Jeopardy!* (and did quite well) in 2011

1. Simple Reflex Agent
2. Model-based Reflex Agent
3. Model-based Goal-based Agent
- ✓ 4. Model-based Utility-based Agent


← Not just to win but to get the most money

# Search

---

# Why Search?

Some problems are small enough to go through all possible states:




Tic-tac-toe:  $3^9 = 19,683$  states (3 values for each cell, nine cells)

(With fewer  
*valid* states)

Most real-world problems are intractable

- We need to be smart about how we visit states




[examples.gurobi.com/traveling-salesman-problem](https://examples.gurobi.com/traveling-salesman-problem)

# Why Search?

---

For when we don't know how to reach a goal (or **any** goal) from a given initial state (or **any** initial state)

# Search Today: AlphaGO



*Pictures from needpix.com, machinelearnings.co/understanding-alphago-948607845bb1*

# Search Today: Story Generation




Figure 1: A narrative search space graph for the plot of *Treasure Island* (Stevenson 1919).

# Defining a Search Problem

---

Break down the problem space into **states** and actions (for each state)

+ these components:

- Initial state – where the agent starts
- Transition model – how actions change states
- Goal test – measuring if the current state is the goal
- Step cost (optional) – how expensive it is to take action  $a$  in state  $s$

Problem: Find a sequence of **actions** that **transitions** the **initial state** into a state which passes the **goal test**

Solution: sequence of actions, or a **plan**

# Formal Definition of a Search Problem

---

1. **States:** a set  $S$

2. **Initial state:**  $s_i \in S$

3. **Actions:** a set  $A$

$\forall s$  **Actions(s)** = the set of actions that can be executed in  $s$ .

4. **Transition Model:**  $\forall s \forall a \in \text{Actions}(s)$

**Result(s, a)**  $\rightarrow s_r$

$s_r$  is called a **successor** of  $s$

$\{s_i\} \cup \text{Successors}(s_i)^* = \text{state space}$

5. **Path cost** (Performance Measure):

Must be additive, e.g. sum of distances, number of actions executed, ...

**c(x,a,y)** is the step cost, assumed  $\geq 0$

◦ (where action **a** goes from state **x** to state **y**)

6. **Goal test:** **Goal(s)**

$s$  is a goal state if **Goal(s)** is true.

Can be implicit, e.g. **checkmate(s)**


# What are States?

---

Current position and attributes of everything in the environment

- Only need the things relevant to the agent's decision making

Like discrete snapshots over time









## State

Location A: vampire, human

Location B: empty

# Vampire World: State Examples

**States:** A snapshot of the world says which entities are in which cells.




And so on...

# What Are Actions?

---

Anything the agent can do to affect the environment




## Actions

*move left, move right, suck*

Or for this particular state...

*move right, suck*

# Vampire World: Transition Model Example




# Goal Test


---

Measuring to see if the goal has been achieved.


If our goal is for the vampire to suck a human's blood, then we need a way of measuring to see if that occurred.



# Sliding Number Puzzle *Search Tree*



# Sliding Number Puzzle *Search Tree*



# Useful Concepts

---

**State space:** the set of all states reachable from the initial state by *any* sequence of actions

- When several operators can apply to each state, this gets large very quickly
- Might be a proper subset of the set of configurations

**Path:** a sequence of actions leading from one state  $s_j$  to another state  $s_k$

**Solution:** a path from the initial state  $s_i$  to a state  $s_f$  that satisfies the goal test

**Search tree:** a way of representing the paths that a search algorithm has explored. The root is the initial state, leaves of the tree are successor states.

**Frontier:** those states that are available for *expanding* (for applying legal actions to)

# Solutions and *Optimal* Solutions

---

A **solution** is a sequence of **actions** from the **initial state** to a **goal state**.

**Optimal Solution:** A solution is **optimal** if no solution has a lower **path cost**.

# Basic search algorithms: Tree Search

---

Generalized algorithm to solve search problems

**Enumerate in some order all possible paths from the initial state**

- Here: search through *explicit tree generation*
  - ROOT= initial state
  - Nodes in search tree generated through **transition model**
  - Tree search treats different paths to the same node as distinct


# Generic Search Algorithm

---

```
function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier
```


# The Generic Search Algorithm

---




# The Generic Search Algorithm

---




# The Generic Search Algorithm

---




# The Generic Search Algorithm

---




# The Generic Search Algorithm

---




# The Generic Search Algorithm

---




# The Generic Search Algorithm

---




# The Generic Search Algorithm

---



# Generalized tree search



function TREE-SEARCH(*problem, strategy*) return a solution or failure

Initialize frontier to the *initial state* of the *problem*  
do

if the frontier is empty then return *failure*

*choose leaf node for expansion according to strategy* & remove from frontier

if node contains goal state then return *solution*

else expand the node and add resulting nodes to the frontier

**The strategy determines search process!**

# The Generic Search Algorithm

**function** TREE-SEARCH(*problem*) **returns** a solution, or failure

    initialize the frontier using the initial state of *problem*

**loop do**

**if** the frontier is empty **then return** failure


        choose a leaf node and remove it from the frontier

**if** the node contains a goal state **then return** the corresponding solution

        expand the chosen node, adding the resulting nodes to the frontier

Challenges:

- Loops
- Inefficiency




# The Generic Search Algorithm

**function** GRAPH-SEARCH(*problem*) **returns** a solution, or failure  
    initialize the frontier using the initial state of *problem*  
    *initialize the explored set to be empty*  
    **loop do**  
        **if** the frontier is empty **then return** failure  
        choose a leaf node and remove it from the frontier  
        **if** the node contains a goal state **then return** the corresponding solution  
        *add the node to the explored set*  
        expand the chosen node, adding the resulting nodes to the frontier  
            *only if not in the frontier or explored set*

Challenges:

- ~~Loops~~
- Inefficiency



# Search Strategies

---

Several classic search algorithms differ only by the order of how they expand their search trees

You can implement them by using different **frontier data structures**

- **Depth-first search** = LIFO queue
- **Breadth-first search** = FIFO queue
- **Greedy best-first search** or **A\* search** = Priority queue

# What about Efficiency?

---

```
function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier
```

Use data structures to control which nodes get explored next!

**How we store the frontier affects the performance of the algorithm.**

# Breadth-First Search (BFS)

---

**function** TREE-SEARCH(*problem*) **returns** a solution, or failure

initialize the frontier using the initial state of *problem*

**loop do**

if the frontier is empty **then return** failure

choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

expand the chosen node, adding the resulting nodes to the frontier


**Stores frontier as first-in first-out (FIFO) queue**

# Breadth-First Search (BFS): FIFO Queue

---

Frontier:  
[Start]

Current node:  
*none*




# Breadth-First Search (BFS): FIFO Queue

---

Frontier:  
[]

Current node:  
Start




# Breadth-First Search (BFS): FIFO Queue

---

Frontier:  
[A, B]

Current node:  
Start




# Breadth-First Search (BFS): FIFO Queue

---

Frontier:  
[B]

Current node:  
A




# Breadth-First Search (BFS): FIFO Queue

---

Frontier:  
[B, C, D]

Current node:  
A




# Breadth-First Search (BFS): FIFO Queue

---

Frontier:  
[C, D]

Current node:  
B




# Breadth-First Search (BFS): FIFO Queue

---

Frontier:  
[C, D, Goal, F]

Current node:  
B



# Breadth-First Search (BFS): FIFO Queue

Frontier:



[C, D, Goal, F]

Eventually...

Current node:  
B

Return:

[Start, B, Goal]



# Properties of Breadth-First Search (BFS)

---

**Complete?**

Yes (if  $b$  is finite)

**Optimal?**

Yes, if cost = 1 per step  
(not optimal in general)

**Time Complexity?**

$1+b+b^2+b^3+\dots+b^d = O(b^d)$

**Space Complexity?**

$O(b^d)$  (keeps every node in memory)

Time/space complexity variables

- $b$ , **maximum branching factor** of search tree
- $d$ , **depth** of the shallowest goal node
- $m$ , **maximum length** of any path in the state space (potentially  $\infty$ )

# Depth-First Search (DFS)

---

```
function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier
```

**Stores frontier as last-in first-out (LIFO) queue, or a stack.**

# BFS versus DFS

---

## Breadth-first

- ✓ Complete,
- ✓ Optimal
- ✗ *but* uses  $O(b^d)$  space

## Depth-first

- ✗ Not complete unless  $m$  is bounded
- ✗ Not optimal
- ✗ Uses  $O(b^m)$  time; terrible if  $m \gg d$
- ✓ *but* only uses  $O(b*m)$  space

## Time/space complexity variables

$b$ , *maximum branching factor* of search tree  
 $d$ , *depth* of the shallowest goal node  
 $m$ , *maximum length* of any path in the state space (potentially  $\infty$ )

# Your Turn – Depth-First Search: LIFO Stack

Frontier:  
[Start]

Current node:  
*none*

Perform DFS on this search tree, listing the frontier and current node at every step until you reach the goal. Please submit to Blackboard when you're done.

